

6.867

Sampling

Fall 2016

Inference is hard



$$P(\mathbf{Z}, \mathbf{W}; \alpha, \beta) = \prod_{j=1}^M \frac{\Gamma\left(\sum_{i=1}^K \alpha_i\right)}{\prod_{i=1}^K \Gamma(\alpha_i)} \frac{\prod_{i=1}^K \Gamma(n_{j,(\cdot)}^i + \alpha_i)}{\Gamma\left(\sum_{i=1}^K n_{j,(\cdot)}^i + \alpha_i\right)} \times \prod_{i=1}^K \frac{\Gamma\left(\sum_{r=1}^V \beta_r\right)}{\prod_{r=1}^V \Gamma(\beta_r)} \frac{\prod_{r=1}^V \Gamma(n_{(\cdot),r}^i + \beta_r)}{\Gamma\left(\sum_{r=1}^V n_{(\cdot),r}^i + \beta_r\right)}.$$

Sampling from $P(Z|W, \alpha, \beta)$ not easy!

LDA model: https://en.wikipedia.org/wiki/Latent_Dirichlet_allocation

Key Points, Topics

Exact inference is hard

Drawing samples from high-dim $P(\mathbf{x})$ hard
even if we can easily evaluate $P^*(\mathbf{x}) \propto P(\mathbf{x})$

Topics

Importance Sampling


MCMC Methods

- Metropolis


- Metropolis-Hastings

- Gibbs Sampling


- Mixing time



Difficulty of estimating $E[f]$




depth at (x,y) is P^*




Task: draw random water samples to estimate plankton conc.

$$\Phi = \frac{1}{Z} \int P^*(x) \phi(x) dx$$


Go to any desired loc, measure P^* and concentration (ϕ) there



Simpler proposal distribution $q(z)$



Gaussian proposal (straight line is true distr)



Cauchy proposal (straight line is true distr)

Metropolis-Hastings

Assume we can evaluate $p^*(z)$

1. Tentative state z' generated from $q(z'|z(t))$

2.
$$a = \frac{p^*(z')}{p^*(z(t))} \frac{q(z^{(t)}; z')}{q(z'; z^{(t)})}$$

3. If $a \geq 1$, accept new state **else** accept with prob. a

4. If accepted, $z(t+1)=z'$ **else** $z(t+1)=z(t)$

For any positive q , i.e., $q(z;z') > 0$, as $t \rightarrow \infty$ prob. distribution of $z(t)$ tends to $p^*(z)/Z$

Gibbs Sampling

Assume: Sampling from conditional is “easy”;
operate as (cyclic) coordinate descent

$$z_1^{(t+1)} \sim p(z_1 | z_2^{(t)}, z_3^{(t)}, \dots, z_n^{(t)})$$

$$z_2^{(t+1)} \sim p(z_2 | z_1^{(t+1)}, z_3^{(t)}, \dots, z_n^{(t)})$$

$$z_n^{(t+1)} \sim p(z_n | z_1^{(t+1)}, z_2^{(t+1)}, \dots, z_{n-1}^{(t+1)})$$

Does it work? Why?

Markov Chains

Initial state: $p^{(0)}(z)$

Transition probability: $T(z'; z)$

Update: $p^{(t+1)}(z') = \int T(z'; z)p^{(t)}(z)dz$

Required properties:

$$p^*(z') = \int T(z'; z)p^*(z)dz \quad \text{(Stationarity)}$$

$$p^{(t)}(z) \rightarrow p^*(z'), \quad \text{as } t \rightarrow \infty \quad \text{(Ergodicity)}$$









A =

$$P(x) = \begin{cases} 1/21 & x \in \{0, 1, 2, \dots, 20\} \\ 0 & \text{otherwise.} \end{cases}$$

$$P(x) = \begin{cases} 1/21 & x \in \{0, 1, 2, \dots, 20\} \\ 0 & \text{otherwise.} \end{cases}$$

Proposal

$$Q(x'; x) = \begin{cases} 1/2 & x' = x \pm 1 \\ 0 & \text{otherwise.} \end{cases}$$



Credit: MacKay, 2005

What's the stationary distribution?

$$A = \begin{pmatrix} & & & 1 \\ & & \frac{1}{2} & \frac{1}{2} \\ & \ddots & & \vdots \\ \frac{1}{n} & \ddots & \ddots & \frac{1}{n} \end{pmatrix}$$

Useful links

MCMC without the bs

Chapter 29, of David MacKay's book:
"Information Theory, Inference, and Learning Algorithms"

Andrieu, de Freitas, Doucet, Jordan.
"Introduction to MCMC for Machine Learning"

See viz. at: <http://setosa.io/ev/markov-chains/>